

NBA Salary Exploratory Data Analysis

Brian Fernando

Aarush Maddela

Nixon Tan

Sharona Yang

Wednesday 17th December, 2025



In this notebook, we'll do some EDA on the dataset containing NBA salary and stats data.

0.1 Importing the data

```
import os
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns

data_path_2024 = os.getcwd() + '/data/final_2024_player.csv'
data_path_2025 = os.getcwd() + '/data/final_2025_player.csv'

df_2024 = pd.read_csv(data_path_2024)
df_2025 = pd.read_csv(data_path_2025)

df = pd.concat([df_2024, df_2025])

df.head()
```

Rk	Player	Age	Team	Pos	G	GS	MP	FG%	FGA	Second_Diff	Third_Diff	DefTeam1	DefTeam2	Next_Year_Salary	Next_Year_Guaranteed	24_contract_year	25_contract_year
0	363	A.J.24.0	MIISG	56.0.0	11.01.5	3.5	...	0	0	0	0	1901	76976912069698				
			Green														
1	476	A.J.20.0	ATSF	20.0.0	8.6	0.9	3.1	...	0	0	0	0	371790084250260000	N			
			Grif-														
2	109	A.Jar.28.	DERF	73.073.0	31.5.5	9.8	...	0	0	0	0	2220	618763284122N	727			
			Gor-														
3	278	A.Jar.27.	HOPG	78.01.0	16.32.4	5.3	...	0	0	0	0	201970	0.0466966000	N			
			Hol-														
4	136	A.Jar.24.	INDS	72.047.027.74.4	8.8	...	0	0	0	0	563325842570660000	Q	0				
			Ne-														
			smith														

5 rows × 67 columns

0.2 Null Values

```
null = df.isnull().sum()

null_dict = dict(zip(null.index, null.values))

[key for key, value in null_dict.items() if value != 0]

['3P%', 'FT%', 'Awards', '2023 -24_contract_year', '2024 -25_contract_year']
```

In this analysis, we won't be using any of these columns, so we can safely drop them.

```
df = df.drop(['3P%', 'FT%', 'Awards', '2023 -24_contract_year', '2024 -25_contract_year'], axis=1)

df.columns

Index(['Rk_x', 'Player', 'Age', 'Team', 'Pos', 'G', 'GS', 'MP_x', 'FG', 'FGA',
       'FG%', '3P', '3PA', '2P', '2PA', '2P%', 'eFG%', 'FT', 'FTA', 'ORB',
       'DRB', 'TRB', 'AST', 'STL', 'BLK', 'TOV', 'PF', 'PTS', 'PER', 'TS%',
       '3PAr', 'FTr', 'ORB%', 'DRB%', 'TRB%', 'AST%', 'STL%', 'BLK%', 'TOV%',
       'USG%', 'OWS', 'DWS', 'WS', 'WS/48', 'OBPM', 'DBPM', 'BPM', 'VORP',
       'Season', 'Experience', 'NumOfAwards', 'All -Star', 'AwardWinner',
       'FirstTeam', 'SecondTeam', 'ThirdTeam', 'DefTeam1', 'DefTeam2',
       'Salary', 'Guaranteed', 'Next_Year_Salary', 'Next_Year_Guaranteed'],
      dtype='object')
```

1 Analysis

1.1 NBA Salary Distribution


```
bins = np.arange(0, df['Next_Year_Salary'].max()+5000000, 5000000)

plt.hist(df['Next_Year_Salary'], bins = bins, edgecolor = 'black')
ax = plt.gca()

ax.xaxis.set_major_formatter(
    plt.FuncFormatter(lambda x, _: f"${x/1e6:.0f}M")
)

plt.xlabel('Next Year Salary')
plt.ylabel('Frequency')
plt.title('Distribution of NBA Player Salaries')

plt.savefig("outputs/salary_distribution.png")
```



The distribution of NBA salaries is clearly right skewed. Most players earn below 5 million dollars. This makes sense, because there are very few “superstars” who are deserving of a salary above 30 million, and majority of NBA players are role players or depth pieces that aren’t expected to make much relative to others in their careers.

1.2 Salary and Stats (BPM, VORP)


```
plt.scatter(df['BPM'], df['Next_Year_Salary'])
ax = plt.gca()
```

```

        ax.yaxis.set_major_formatter(
            plt.FuncFormatter(lambda x, _: f"${x/1e6:.0f}M")
        )
        plt.xlabel("Box Plus Minus")
        plt.ylabel("Next Year Salary")
        plt.title("Box Plus Minus Vs. Salary")

Text(0.5, 1.0, 'Box Plus Minus Vs. Salary')

```



Box plus minus is a popular basketball advanced statistic that estimates a player's contribution to the team's performance, both offensively and defensively. A higher box plus minus score means that a player is contributing more towards their team's success. It seems like there is a general positive trend in that players with higher BPM scores can expect to see higher salaries. However, there are a couple outliers, so we'll subset the dataset to only include players with BPM scores higher than -10 and replot this relationship.

```


subset = df[df['BPM'] > -10]

plt.scatter(subset['BPM'], subset['Next_Year_Salary'])
ax = plt.gca()

ax.yaxis.set_major_formatter(
    plt.FuncFormatter(lambda x, _: f"${x/1e6:.0f}M")
)
plt.xlabel("Box Plus Minus")
plt.ylabel("Next Year Salary")
plt.title("Box Plus Minus Vs. Salary")

Text(0.5, 1.0, 'Box Plus Minus Vs. Salary')

```



```

print(f"Correlation coefficient: {df['BPM'].corr(df['Next_Year_Salary'])}")

print(f"Subsetted Correlation coefficient: {subset['BPM'].corr(subset['Next_Year_Salary'])}")

Correlation coefficient: 0.5565013579110638
Subsetted Correlation coefficient: 0.5871708693441092

```

After removing the outliers, the scatterplot gives a better view of the relationship between BPM and NBA salaries. The correlation coefficient confirms that there is a moderately strong positive association between the two.

Let's do the same with VORP now


```

plt.scatter(df['VORP'], df['Next_Year_Salary'])
ax = plt.gca()

ax.yaxis.set_major_formatter(
    plt.FuncFormatter(lambda x, _: f"${x/1e6:.0f}M")
)
plt.xlabel("VORP Minus")
plt.ylabel("Next Year Salary")
plt.title("VORP Vs. Salary")

Text(0.5, 1.0, 'VORP Vs. Salary')

```



```
print(f"Correlation coefficient: {df['VORP'].corr(df['Next_Year_Salary'])}")

Correlation coefficient: 0.6744527184663939
```

VORP is another popular advanced stat that estimates a player's overall contribution to a team compared to an average NBA player. It shows the same positive association with NBA salaries as BPM did; however, the correlation coefficient does show a stronger linear relationship than with BPM.

I'll combine both plots into one figure

```
fig, axs = plt.subplots(nrows = 1, ncols = 2, figsize = (10,6))

plt.subplots_adjust(wspace=0.35)

axs[0].scatter(subset['BPM'], subset['Next_Year_Salary'])
axs[1].scatter(df['VORP'], df['Next_Year_Salary'])


axs[0].set_xlabel("Box Plus Minus (BPM)")
axs[0].set_ylabel("Salary")
axs[0].set_title("Box Plus Minus Vs. Salary")

axs[1].set_xlabel("Value Over Replacement Player (VORP)")
axs[1].set_ylabel("Salary")
axs[1].set_title("VORP Vs. Salary")

formatter = plt.FuncFormatter(lambda y, _: f"${y/1e6:.0f}M")

for ax in axs:
    ax.yaxis.set_major_formatter(formatter)

plt.savefig("outputs/bpm_vorp_vs_salary.png")
```



1.3 Offensive Metrics vs. Defensive Metrics and Their Impacts on Salaries

Let's take a look at how offensive and defensive metrics relate to NBA salaries

```
good_scoring = df[(df['FGA'] >= 10) & (df['FG%'] >= 0.45)]
bad_scoring = df[~(df['FGA'] >= 10) & (df['FG%'] >= 0.45)]

plt.boxplot([good_scoring['Next_Year_Salary'], bad_scoring['Next_Year_Salary']], labels = [ 


ax = plt.gca()

ax.yaxis.set_major_formatter(
    plt.FuncFormatter(lambda x, _: f"${x/1e6:.0f}M")
)

plt.ylabel('Salary')
plt.title('Salary Comparison Between Prolific Scoring Players and Non Prolific Scoring Players')

plt.savefig("outputs/scoring_volume_salary_comparison.png")

/tmp/ipykernel_1497/3721848652.py:4: MatplotlibDeprecationWarning: The 'labels' parameter of
```



This plot shows that there is quite a big difference in salaries between efficient scorers in the league and those who aren't. Moreover, the boxplot for prolific scorers is about normally distributed, while the boxplot for nonprolific scorers is definitely right skewed. This shows that most players who aren't efficiently scoring for their teams earn significantly less than those that do. This is probably a big factor that goes into determining a player's salary, which makes sense, as basketball is primarily an offensively minded sport.

```

fig, axs = plt.subplots(nrows = 1, ncols = 2, figsize=(11,6))

plt.subplots_adjust(wspace=0.35)

axs[0].scatter(df['OBPM'], df['Next_Year_Salary'])
axs[1].scatter(df['DBPM'], df['Next_Year_Salary'])

axs[0].set_xlabel("OBPM")
axs[0].set_ylabel("Salary")
axs[0].set_title("Offensive Box Plus Minus Vs. Salary")


axs[1].set_xlabel("DBPM")
axs[1].set_ylabel("Salary")
axs[1].set_title("Defensive Box Plus Minus Vs. Salary")

formatter = plt.FuncFormatter(lambda y, _: f"${y/1e6:.0f}M")

for ax in axs:
    ax.yaxis.set_major_formatter(formatter)

plt.savefig("outputs/obpm_dbpm_salary_comp.png")

```



```

print(f"OBPM Correlation coefficient: {df['OBPM'].corr(df['Next_Year_Salary'])}")
print(f"DBPM Correlation coefficient: {df['DBPM'].corr(df['Next_Year_Salary'])}")

OBPM Correlation coefficient: 0.6401944026484849
DBPM Correlation coefficient: 0.05995507005883653

```

It's difficult to group players into defensively good and bad categories using defensive stats, as they aren't as accurate as offensive stats in evaluating a player's effectiveness. Instead, we used defensive box plus minus, which is essentially the same as box plus minus, but it evaluates a player's defensive efficiency/potential. We compare this to its counterpart, offensive box plus minus. The scatterplots and correlation coefficients both show that there isn't much of a relationship between defensive BPM and salaries, whereas, there is a moderately strong positive correlation between offensive BPM and salaries.

The last two plots show that offensive potential is more valued than defensive potential in evaluating a player's worth to a team through their salary. There is a clear salary difference between good offensive players and more average players, while there is not much of a relationship between defensive stats and NBA salaries.

```

fig, axs = plt.subplots(nrows = 1, ncols = 2, figsize = (12, 6))

axs[0].scatter(df['OBPM'], df['BPM'])
axs[1].scatter(df['DBPM'], df['BPM'])


axs[0].set_xlabel("Offensive Box Plus Minus")
axs[0].set_ylabel("Box Plus Minus")

axs[1].set_xlabel("Defensive Box Plus Minus")
axs[1].set_ylabel("Box Plus Minus")

axs[0].set_title("OBPM Vs. BPM")
axs[1].set_title("DBPM Vs. BPM")

plt.savefig("outputs/obpm_dbpm_vs_bpm.png")

```



Overall BPM is the sum of Offensive BPM and Defensive BPM. Our earlier analysis shows that there is a higher correlation between offensive stats and salaries than for defensive stats. We also showed a general positive association between BPM and salaries. Out of curiosity, I wanted to see the relationship between OBPM/DBPM and general BPM. There is clearly a stronger positive relationship between OBPM and BPM than for DBPM, which aligns with our earlier analysis between stats and salaries.

From a salary standpoint, it seems like an offensive superstar is worth more than a defensive superstar.


1.4 Salary In Relation to Experience and Awards

```
age_salary = df.groupby('Age')['Next_Year_Salary'].mean()

ax = plt.gca()

ax.yaxis.set_major_formatter(
    plt.FuncFormatter(lambda x, _: f"${x/1e6:.0f}M")
)

plt.bar(age_salary.index, age_salary.values)
plt.xlabel("Age")
plt.ylabel("Avg Salary")
plt.title("Average Salary by Age")
plt.show()
```



```
df[df['Age'] == 40]
```

Rk	Playe	Tea	RosG	GS	MPF	G	FGA.	Awa	Fri	WT	TH	TT	St	G2	Name	Ne	Year	Sal	Guaranteed
20119	(LeBron) James	LASF	70.0	70.0	84.9	3.3	18.1..	0	0	1	0	0	0	487487	12627153				

1 rows × 62 columns


```
subset = df[df['Age'] < 39]
age_salary = subset.groupby('Age')['Next_Year_Salary'].mean()

ax = plt.gca()

ax.yaxis.set_major_formatter(
    plt.FuncFormatter(lambda x, _: f"${x/1e6:.0f}M")
)

plt.bar(age_salary.index, age_salary.values)
plt.xlabel("Age")
plt.ylabel("Avg Salary")
plt.title("Average Salary by Age")

plt.savefig("outputs/avg_salary_by_age.png")
```



The 40 years old column stands out, but after inspecting the dataframe, it makes sense, as there is only one 40 year old player, which is Lebron James. Arguably the greatest player of all time, he has a salary of about \$48 million. Other than that though, the graph looks to be bimodal, with peaks around 29 years and 34 years. This makes sense, as players in their late 20s are typically around their peaks, which would coincide with higher salaries. The second peak is around older players, but the only players to remain in the league in that age range are typically older superstars whose careers are defined by longevity. It then falls off, as very old players in their late 30s aren't as efficient as they used to be, which coincides with smaller salaries.

```

awards_salary = df.groupby('NumOfAwards')['Next_Year_Salary'].mean()

plt.bar(awards_salary.index, awards_salary.values)

ax = plt.gca()


ax.yaxis.set_major_formatter(
    plt.FuncFormatter(lambda x, _: f"${x/1e6:.0f}M")
)

plt.xlabel('Number of Awards')
plt.ylabel("Average Salary")

plt.title("Average Salary By Number of Awards")

plt.savefig("outputs/avg_salary_by_awards.png")

```



```

all_star = df[df['All -Star'] == 1]
non_all_star = df[df['All -Star'] == 0]

plt.boxplot([all_star['Next_Year_Salary'], non_all_star['Next_Year_Salary']], labels = ["All Star", "Non All Star"])
ax = plt.gca()


ax.yaxis.set_major_formatter(
    plt.FuncFormatter(lambda x, _: f"${x/1e6:.0f}M")
)

plt.ylabel("Salary")
plt.title("NBA Salary By All Star Selection")

plt.savefig("outputs/all_star_salary_comp.png")

/tmp/ipykernel_1497/133333447.py:4: MatplotlibDeprecationWarning: The 'labels' parameter of
plt.boxplot([all_star['Next_Year_Salary'], non_all_star['Next_Year_Salary']], labels = ["All Star", "Non All Star"])

```



There are many awards in the NBA most notably: Most Valuable Player (MVP), Defensive Player of the Year (DPOY), and All NBA team selections. As expected, players with more awards tend to have higher salaries, as those with more awards are naturally better players. Another important accolade is selection for the All Star team. The side by side boxplot shows that All Star Players do earn much more than those who aren't. However, there are a lot of non All Star players as outliers who make about the same as All Star players. This is because All Star nominations are extremely selective, as there is only a certain amount of players each year who can be selected. As a result, there are many extremely talented players who are left off of the All Star roster.

1.5 Salary By Position

```
position_salary = df.groupby('Pos')['Next_Year_Salary'].mean()

plt.bar(position_salary.index, position_salary.values)


ax = plt.gca()

ax.yaxis.set_major_formatter(
    plt.FuncFormatter(lambda x, _: f"${x/1e6:.0f}M")
)

plt.xlabel("Position")
plt.ylabel("Average Salary")

plt.title("Average Salary By Position")

plt.savefig("outputs/salary_by_pos.png")
```



It seems like point guards on average are paid more than their peers. This makes sense, because usually the point guards are the players managing a team's offense. Oftentimes, team offenses are largely structured around the play of their pointguards. Due to their importance to teams, it does make sense that they are paid more on average than other positions. All other positions are on average paid around the same at around \$12 million.

2 Multicollinearity and VIF Analysis

NBA Stats data is highly correlated, as most advanced stats are functions of other stats. Any offensive percentage based statistic is simply the quotient of successful attempts and total attempts. For instance, $FG\% = FG/FGA$ (Field Goals Made / Field Goals Attempted). It's important to conduct a VIF analysis and remove highly correlated variables, as they can cause instability in our later regression models.

2.1 Current VIF Scores

```
vif_columns = ['Age', 'G', 'GS', 'MP_x', 'FG', 'FGA',
               'FG%', '3P', '3PA', '2P', '2PA', '2P%', 'eFG%', 'FT', 'FTA',
               'ORB', 'DRB', 'TRB', 'AST', 'STL', 'BLK', 'TOV', 'PF', 'PTS', 'PER', 'TS%', '3PAr',
               'STL%', 'BLK%', 'TOV%', 'USG%', 'OWS', 'DWS', 'WS', 'WS/48', 'OBPM',
               'DBPM', 'BPM', 'VORP', 'Experience', 'NumOfAwards',
               'All -Star', 'AwardWinner', 'FirstTeam', 'SecondTeam', 'ThirdTeam',
               'DefTeam1', 'DefTeam2']

non_numeric = ['Player', 'Team', 'Pos']

from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.preprocessing import StandardScaler
```

```
df_vif = df[vif_columns].copy()
df_vif = df_vif.fillna(df_vif.median())

scaler = StandardScaler()
X_scaled = scaler.fit_transform(df_vif)

vif_df = pd.DataFrame()
vif_df['feature'] = df_vif.columns
vif_df['VIF'] = [variance_inflation_factor(X_scaled, i)
                 for i in range(X_scaled.shape[1])]

vif_df.sort_values('VIF', ascending = False)
```

	feature	VIF	
5	FGA	11314.221656	
23	PTS	9776.988257	
4	FG	6132.743387	
10	2PA	5944.950433	
42	BPM	5438.085530	
40	OBPM	3779.923247	
38	WS	3567.776088	
8	3PA	2562.969471	
17	TRB	2562.891993	
9	2P	2021.851050	
36	OWS	1787.044350	
16	DRB	1476.600673	
30	TRB%	1404.880118	
41	DBPM	895.620263	
7	3P	738.359541	
29	DRB%	627.270609	
13	FT	535.581498	
37	DWS	500.505318	
6	FG%	415.821142	
12	eFG%	303.521027	
28	ORB%	278.749742	
15	ORB	270.179573	
24	PER	225.317081	
14	FTA	141.136785	
26	3PAr	101.853601	
39	WS/48	101.298727	
25	TS%	79.766231	
3	MP_x	75.411293	
35	USG%	54.571440	
43	VORP	46.427325	
18	AST	44.591487	
21	TOV	30.921930	
31	AST%	25.900951	
33	BLK%	14.185062	
20	BLK	14.082939	
19	STL	11.970258	
32	STL%	11.292303	
45	NumOfAwards	9.918648	
44	Experience	9.062526	
0	Age	8.625368	
34	TOV%	7.526962	
11	2P%	7.413079	
27	FTr	6.024722	
1	G	5.453902	
2	GS	4.688047	
22	PF	17	4.434805
48	FirstTeam	3.847879	
46	All-Star	3.490438	
49	SecondTeam	2.462567	
51	DefTeam1	1.943521	
52	DefTeam2	1.653303	
50	ThirdTeam	1.638871	
47	AwardWinner	1.386805	

There are a lot of highly collinear variables, as shown by the high VIF scores. Let's drop all of the counting stats and keep percentage based stats. We'll also drop most of the advanced stats but keep BPM, as most advanced stats are factored into BPM.

2.2 Dropping Collinear Columns

```
cols2drop = ['G', 'GS', 'FG', 'FGA', '3P', '3PA', '2P', '2PA', 'FT', 'FTA', 'ORB', 'DRB', 'T']

df_vif.drop(cols2drop, axis=1, inplace=True)

df_vif.columns

Index(['Age', 'MP_x', 'PF', 'TS%', 'TRB%', 'AST%', 'STL%', 'BLK%', 'TOV%',
       'USG%', 'BPM', 'NumOfAwards', 'All -Star', 'AwardWinner', 'FirstTeam',
       'SecondTeam', 'ThirdTeam', 'DefTeam1', 'DefTeam2'],
      dtype='object')

df_vif = df_vif.fillna(df_vif.median())

scaler = StandardScaler()
X_scaled = scaler.fit_transform(df_vif)

vif_df = pd.DataFrame()
vif_df['feature'] = df_vif.columns
vif_df['VIF'] = [variance_inflation_factor(X_scaled, i)
                 for i in range(X_scaled.shape[1])]

vif_df.sort_values('VIF', ascending = False)
```

	feature	VIF
11	NumOfAwards	9.163318
10	BPM	9.101377
5	AST%	4.024821
1	MP_x	3.598090
3	TS%	3.450686
14	FirstTeam	3.128119
12	All-Star	3.072324
8	TOV%	2.705655
2	PF	2.578249
15	SecondTeam	2.320510
4	TRB%	2.306702
7	BLK%	2.070009
9	USG%	2.003561
17	DefTeam1	1.746462
6	STL%	1.717763
18	DefTeam2	1.557843
16	ThirdTeam	1.513444
13	AwardWinner	1.327046
0	Age	1.118065

Much better, we can keep all of these columns, as the VIF values are all less than 10.

```
cols2keep = list(vif_df.feature)

cols2keep

['Age',
 'MP_x',
 'PF',
 'TS%',
 'TRB%',
 'AST%',
 'STL%',
 'BLK%',
 'TOV%',
 'USG%',
 'BPM',
 'NumOfAwards',
 'All -Star',
 'AwardWinner',
 'FirstTeam',
 'SecondTeam',
 'ThirdTeam',
 'DefTeam1',
```

```

'DefTeam2']

final_df = df[cols2keep+['Pos', 'Next_Year_Salary']]

final_df.columns

Index(['Age', 'MP_x', 'PF', 'TS%', 'TRB%', 'AST%', 'STL%', 'BLK%', 'TOV%',
       'USG%', 'BPM', 'NumOfAwards', 'All -Star', 'AwardWinner', 'FirstTeam',
       'SecondTeam', 'ThirdTeam', 'DefTeam1', 'DefTeam2', 'Pos',
       'Next_Year_Salary'],
      dtype='object')

display(final_df)

```

	Age	MP_x	PF	TS%	TRB%	AST%	STL%	BLK%	TOV%	USG%	BPM	NumOfAwards	All -Star	AwardWinner	FirstTeam	SecondTeam	ThirdTeam	DefTeam1	DefTeam2	Pos	Next_Year_Salary
Star																					
0	24.0	1.0	0.9	0.6	15.7	6.4	0.7	0.6	5.5	15.2..	0	0	0	0	0	0	0	0	0	SG 2120693	
1	20.0	8.6	0.3	0.3	27.3	7.3	0.3	1.1	11.3	17.3..	0	0	0	0	0	0	0	0	0	SF 250000	
2	28.0	1.5	1.9	0.6	0.7	11.4	5.2	1.2	1.7	11.2	7.8..	0	0	0	0	0	0	0	0	PF 22841455	
3	27.0	1.6	3.1	0.6	0.5	22.8	2.2	15.4	1.6	0.4	10.6	16.7..	0	0	0	0	0	0	0	PG 4668000	
4	24.0	27.3	3.3	0.6	31.8	6.6	1.5	2.0	8.4	16.0..	0	0	0	0	0	0	0	0	SF 11000000		
...	
30	727.0	15.3	2.1	0.6	0.5	15.6	1.4	2.6	15.1	7.1..	0	0	0	0	0	0	0	0	0	PF 18080496	
308	22.0	21.5	2.8	0.6	22.0	3.6	1.1	2.5	5.5	15.3	6.4..	1	0	0	0	0	0	0	0	C 6045000	
309	29.0	85.2	1.6	0.6	30.5	18.3	1.1	0.4	13.4	25.3..	0	0	0	0	0	0	0	0	0	SG 47499660	
310	24.0	10.7	1.1	0.5	78.1	4.6	2.0	5.6	8.0	12.2..	0	0	0	0	0	0	0	0	0	PF 8177778	
311	24.0	28.0	2.7	0.6	0.0	36.6	0.2	1	3.0	12.8	4.7..	0	0	0	0	0	0	0	0	PF 39446090	

751 rows × 21 columns

```

final_df.info()

<class 'pandas.core.frame.DataFrame'>
Index: 751 entries, 0 to 311
Data columns (total 21 columns):
 #   Column            Non-Null Count  Dtype  
 --- 
 0   Age               751 non-null    float64
 1   MP_x              751 non-null    float64
 2   PF                751 non-null    float64
 3   TS%               751 non-null    float64
 4   TRB%              751 non-null    float64
 5   AST%              751 non-null    float64
 6   STL%              751 non-null    float64
 7   BLK%              751 non-null    float64

```

```

8   TOV%           751 non -null    float64
9   USG%           751 non -null    float64
10  BPM            751 non -null    float64
11  NumOfAwards   751 non -null    int64
12  All_Star       751 non -null    int64
13  AwardWinner   751 non -null    int64
14  FirstTeam     751 non -null    int64
15  SecondTeam    751 non -null    int64
16  ThirdTeam     751 non -null    int64
17  DefTeam1      751 non -null    int64
18  DefTeam2      751 non -null    int64
19  Pos            751 non -null    object
20  Next_Year_Salary 751 non -null    int64
dtypes: float64(11), int64(9), object(1)
memory usage: 129.1+ KB

```

Above, we printed out the type of each column/feature. We see that all columns are of type int or float except the ‘Pos’ column which is of type of object currently. We want to transform this column from type object to bool using one-hot encoding and then cast the bool values into integer values.

```

oh_df = pd.get_dummies(final_df, columns=['Pos'])
oh_df['Pos_C'] = oh_df['Pos_C'].astype(int)
oh_df['Pos_PF'] = oh_df['Pos_PF'].astype(int)
oh_df['Pos_PG'] = oh_df['Pos_PG'].astype(int)
oh_df['Pos_SF'] = oh_df['Pos_SF'].astype(int)
oh_df['Pos_SG'] = oh_df['Pos_SG'].astype(int)

display(oh_df)

```

	Age	MPPR	TS%	TRAS%	TRBL%	USG%	SecondTeam	ThirdTeam	DefTeam1	DefTeam2	Pos_C	Pos_PF	Pos_PG	Pos_SF	Pos_SG
0	24.0	1.0	0.9	0.6	17.7	6.4	0.7	0.6	5.5	15.2..	0	0	0	0	2120693
1	20.0	8.6	0.3	0.3	27.3	7.0	0.3	1.1	11.3	17.3..	0	0	0	0	2500000
2	28.0	1.5	1.9	0.6	07.1	415.2	1.2	1.7	11.2	17.8..	0	0	0	0	22801455
3	27.0	16.3	1.6	0.5	32	15.4	1.6	0.4	10.6	6.7..	0	0	0	0	4668000
4	24.0	27.7	3.3	0.6	318	6.6	1.5	2.0	8.4	16.0..	0	0	0	0	11000000
...
30727	0	15.3	2.1	0.6	05.6	15.6	1.4	2.6	15.1	7.1..	0	0	0	0	18080496
30822	0	21.5	2.8	0.6	220.5	6.1	1.2	5.5	15.3	6.4..	0	0	0	0	6045000
30929	0	35.2	1.6	0.6	39.5	18.3	1.1	0.4	13.4	25.3..	0	0	0	0	47499660
31024	0	10.7	1.1	0.5	7841	4.6	2.0	5.6	8.0	12.2..	0	0	0	0	8170778
31124	0	28.6	2.7	0.6	03.6	6.6	0.1	3.0	12.8	4.7..	0	0	0	0	39406090

751 rows × 25 columns

```

oh_df.info()

<class 'pandas.core.frame.DataFrame'>
Index: 751 entries, 0 to 311
Data columns (total 25 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Age              751 non-null    float64
 1   MP_x             751 non-null    float64
 2   PF               751 non-null    float64
 3   TS%              751 non-null    float64
 4   TRB%             751 non-null    float64
 5   AST%             751 non-null    float64
 6   STL%             751 non-null    float64
 7   BLK%             751 non-null    float64
 8   TOV%             751 non-null    float64
 9   USG%             751 non-null    float64
 10  BPM              751 non-null    float64
 11  NumOfAwards      751 non-null    int64  
 12  All_Star          751 non-null    int64  
 13  AwardWinner       751 non-null    int64  
 14  FirstTeam          751 non-null    int64  
 15  SecondTeam         751 non-null    int64  
 16  ThirdTeam          751 non-null    int64  
 17  DefTeam1           751 non-null    int64  
 18  DefTeam2           751 non-null    int64  
 19  Next_Year_Salary  751 non-null    int64  
 20  Pos_C              751 non-null    int64  
 21  Pos_PF             751 non-null    int64  
 22  Pos_PG             751 non-null    int64  
 23  Pos_SF             751 non-null    int64  
 24  Pos_SG             751 non-null    int64  
dtypes: float64(11), int64(14)
memory usage: 152.5 KB

```

After one-hot encoding, we see that the original ‘Pos’ column was dropped and replaced by 5 new columns, each corresponding to a ‘Pos’ value and we are now ready to fit models on this processed dataset.

```
oh_df.to_csv("data/oh_df.csv", index=False)
```